

ELC 3338 Project Book

Steve Potter

December 22, 2021

Contents

1	Program Counter Register	4
1.1	Program Counter Register	4
1.2	Testbench	7
1.3	Your Assignment	8
2	Program Counter Incrementer and Mux	10
2.1	Adder	10
2.2	Adder Test Bench	11
2.3	Input Selection via Mux	11
2.4	Your Assignment	13
3	Fetch Stage	15
3.1	Instruction Memory Module	15
3.2	Fetch Stage	16
3.3	Your Assignment	17
4	Beginning to Decode	19
4.1	Instruction Decode	19
4.2	Register File	20
4.3	Your Assignment	23
5	Control Unit and Sign Extender	24
5.1	Control Unit	24
5.2	Control Unit Test	25
5.3	Sign Extender	26
5.4	Sign Extender Test	26
5.5	Your Assignment	27
6	Finishing Decode	28
6.1	iDecode Module	28
6.2	iDecode Testbench	28
6.3	Your Assignment	30
7	Integrating Fetch and Decode	32
7.1	Your Assignment	32

8 ALU and ALU Control	34
8.1 ALU	34
8.2 ALU Control	35
8.3 Your Assignment	36
9 Execute Stage	38
9.1 Execute	38
9.2 Your Assignment	40
10 Integrating Fetch and Decode	41
10.1 Integration	41
10.2 Your Assignment	42
11 Memory	44
11.1 Memory Stage	45
11.2 Data Memory	45
11.3 Branch Resolution	45
11.4 Test Bench	45
11.5 Your Assignment	46
12 Write Back	47
12.1 Mux	47
12.2 Datapath	47
12.3 Your Assignment Part 1	47
12.4 Division	50
12.5 Your Assignment Part 2	52
13 Pipeline Fetch and Decode	53
13.1 Pipeline Fetch and Decode	53
13.2 Pipeline Buffering and Forwarding	53
13.3 Your Assignment	54
14 Pipelining without Branching or Forwarding	55
14.1 Overview	55
14.2 Your Assignment	56

Lab 1

Program Counter Register

During the course of this semester, we will build a 64-bit computer. To do this, we will make a synthesizable machine in Verilog, a common hardware description language (HDL).

A computer runs a program by executing individual instructions in sequential order. The instructions are stored in memory and are accessed by their memory address. During each clock cycle, an instruction is fetched from memory and executed on the processor. The memory address of the next instruction to fetch is stored in a register called the Program Counter (PC). During Lab 1, we will build and test the Program Counter register. In Lab 2, we build an incrementer (to count to the next instruction) and a mux (to select between the incremented count or a new starting value).

1.1 Program Counter Register

In order to make the Program Counter, we are going to make a Verilog module that explains how to build a register (a D flip-flop). Let me unpack the previous sentence:

1. Verilog is a Hardware Description Language (HDL).
2. We write Verilog code to tell Vivado how we want our register module to behave.
3. Vivado reads our Verilog code and synthesizes a realizable digital hardware design that meets the behavior that we specified. Thank you Vivado!
4. Vivado also simulates the behavior of the hardware, allowing you to test your design without building/programming hardware.

Consider the Verilog code in Listing 1.1. It is made up of three sections:

1. Header - this code includes a file name definitions.vh which contains information necessary for the program to run.

2. Port list (also known as interface) - specifies the signals coming into or going out of the module. In this case, there are three inputs and one output.
3. Body (also known as implementation) - describes the functionality of the module.

Listing 1.1: Verilog code to make a register.

```

'include "definitions.vh"

module register(
    input clk ,
    input reset ,
    input  ['WORD-1:0] D,
    output reg ['WORD-1:0] Q
);
    always @(posedge( clk ),posedge( reset )) begin
        if ( reset==1'b1 )
            Q<=WORD'b0;
        else
            Q <= D;
    end
endmodule

```

The first part is the header. We will use this same header each time. It tells the Verilog compiler to get all the data from a file called definitions.vh. The extension vh is a Verilog header. We use this to specify common pieces of data we will use across our design, so that all the components we build will be consistent. By putting them in one file, we make it easier to maintain, and prevent mistakes that can happen easily by having multiple copies of these basic pieces of data. For our first lab, one item that we will be using from definitions.vh is WORD (set to 64), which is the size (in bits) of the memory addresses in our computer. Note that if we build things based around WORD, rather than the number 64, we can just change the value of WORD in the file and get a computer with a different size with a couple key strokes.

The second part is the port list or interface. In this area we specify what signals are coming in (input), going out (output), or could go either direction (inout). Ports can be defined as either "wire" or "reg". This can be confusing to some students. Think of it this way:

1. Wire

- (a) A wire is just a conductor that connects one component or module to another.
- (b) The value on a wire can only be changed by using combinational logic (as opposed to sequential logic).

Figure 1.1: Module Diagram.



- Internal regs are set in the 'initial' block or an 'always' block

- (c) It has no memory, meaning that the value on the wire is driven by the results of combinational logic at that particular moment.
- (d) Module inputs are always wires.
- (e) Module outputs can be wires or regs.

2. Reg

- (a) A reg more closely resembles a variable in software programming languages.
- (b) A value of a reg can only be set by using sequential logic.
- (c) A reg has memory, meaning that the value of the reg will remain the same until a sequential logic element updates it.
- (d) You can directly set a reg to a value using a procedural assignment.
- (e) Regs can be used internally in a module (neither input nor output), or they can be used as module outputs. They cannot be used as module inputs.

If you don't specify anything for the port type, you will get a wire - it is the default. In our case we have four signals: three inputs (always wires), and one output that is a register. The first two inputs are single-bit wires. One is the clock, which specifies the timing, and the other is reset, which clears the contents of the register (makes them zero). The final input is the value we want to store in memory, and I have called it D, following the convention of digital logic. D has multiple bits that are numbered from WORD-1 down to 0. Thus the leftmost bit is 63 in this case, and the rightmost bit is 0¹. The output Q

¹If you want to be technical this is called little endian, since the little end (the least significant or unit bit) is going into the first memory location (bit 0). If you reversed the order by putting the 0 first and the WORD-1 last it would be big endian, since the big end (most significant bit) would go in the lowest addressed bit.

(also the digital logic conventional name) is a register (it will hold its value) and should also be of size WORD and follow the same bit order as the input D.

To help clarify this, please examine Figure 1.1, which shows the interconnection of the modules in this lab.

The final section is the body or implementation. It is composed of a single thread of code, that will keep running (hence always). It will run one time every time there is a positive edge (0 to 1 transition) for either the clock or reset. Reset has higher priority, so if reset is asserted the register is cleared (Q is set to zero), otherwise the value of D is stored in Q. That is it. A nice, simple module. Please note that the provided register module is fully operational. You do not need to modify it. This file is provided for you in ARM-Lab/code/1_fetch/register.v.

1.2 Testbench

We now want to test our register module using System Verilog. System Verilog is very similar to Verilog, but it adds the ability to verify that we get the results that we are expecting. To test the register module, we need to tell the simulator to build a copy (instantiate) of the register module, and then we will need to supply the inputs and evaluate the outputs to verify that the module works correctly. Evaluate the testbench, ARM-Lab/code/1_fetch/register_test.sv.

When evaluating this testbench module (register_test), notice that there are no ports. A testbench is providing all the signals to simulate the inputs to the unit under test (UUT) and thus does not need ports. This is how Verilog finds a top level simulation module - there are no ports. The clock signal will be driven by a module named oscillator, which will give us a square wave with period CYCLE, which is another constant defined in our definitions.vh file. The code thus makes an oscillator and a register, then runs the 'initial' section (it runs once at the start then never again).

This initial section of the testbench follows a relatively simple pattern:

1. The inputs to the register module are set to particular values.
2. The system delays for some amount of time. For instance, a one cycle delay (10ns) is inserted with #CYCLE. When analyzing the system, please note that the clock signal has rising edges at t=5,15,25,35,etc.
3. The value of er (expected result) is set to the output value that you expect to get from the register module.
4. The value and size of er are compared with the value and size of ar (actual result). The actual result is the output of your register module. I provide (in "verification_functions.sv", included at the top of the testbench) the verify function and a few other functions that allow us to easily verify the behavior of our system. Each time the 'verify' function is called, it reports and keeps track of whether the test passed (ar == er) or failed (ar != er).

5. At the end of the testbench, the `final_result` function is called to report the results of the test. This function will show the number of passing and failing test cases.

1.3 Your Assignment

You are to:

1. Evaluate the provided testbench. It is not an exhaustive testbench, but it tests a number of cases that commonly occur in our system. Note that you should not change the input values or timing of the testbench, nor should you add additional test cases to the testbench.
2. Create an Expected Results Table for your testbench. The idea behind the Expected Results Table is that you should identify how you think the system should operate. If you don't know how it should work, you will not know whether your simulation results are correct. The Expected Results Table should have a row for each signal in your simulation results (and the row order should match between your Expected Results Table and Simulation Results). The table should also have a column for each test case in the testbench. These test cases are the points in time that correspond to the 'verify' function calls in the testbench.

I have provided an Expected Results Table for Lab 1 in ARM-Lab/tables. It includes everything you need, except that the row for `er_q` is not filled in. Please fill in this row with the expected value of `q` for each test case. Note that you don't need to show the `clk` signal in the Expected Results Table.

3. The provided testbench does not set `er_q` to the correct value, therefore causing your test to fail. Your job now is to take the values from your expected results table and enter these values in the testbench as `er_q` (expected result of `q`).
4. Run a behavioral simulation. Evaluate the timing diagram and verify that it matches the Expected Results Table. Also evaluate the printouts in the Tcl Console window in Vivado. These printouts will indicate the number of passes and fails that occurred in the test. If you chose the correct `er` values and all tests pass, then your module is verified to work properly.
5. Produce a landscape mode PDF called `Lab1_lastname.pdf` that includes (in this order):
 - (a) Your name and the lab number.
 - (b) A snip (using the Snipping Tool) of your Expected Results Table.
 - (c) A snip of the Simulation Results (the black window that shows each signal). Make sure to show all values in decimal form and don't cut off the signal names on the left.

- (d) The test results copied and pasted from the Tcl Console. The results should show the entire log from BEGIN TEST RESULTS to END TEST RESULTS.
 - (e) I have included a sample, ARM-Lab/reports/Lab1.Potter.pdf to give you an idea of what I'm looking for.
6. Upload Lab1_lastname.pdf file to Canvas.
 7. Zip up your ARM-Lab directory and submit it on Canvas as well. Please make sure that you give me the ARM-Lab directory rather than the ARM-Project directory. I do not want the project files in the ARM-Project directory. Before you submit your zip file, make sure that it contains directories like code, manual, tables, etc. I will extract your zip file and run your code against my correct testbench to verify that your code and testbench work correctly. Since I give you working register.v code in this lab, this is pretty easy. In future labs, you must create your own module code.

Lab 2

Program Counter Incrementer and Mux

As mentioned in the last lab, the program counter is a register that is one word in length. It holds the address in memory of the next instruction to be fetched and executed. There are several ways that the program counter is updated:


1. If the program does not branch (via an if statement, while loop, etc), then the program counter should advance to the next address (by adding 4 to the current PC) each clock cycle.
2. If the conditions of a conditional branch are met, then the program counter should be updated with the branch destination address.
3. If an unconditional branch or jump occurs, then the program counter should be updated with the branch destination address.
4. If an interrupt or error occurs, then the program counter should be updated with the interrupt or error handler address.

The instructions will be fetched in sequential order the majority of the time.

2.1 Adder

We will build a program counter incrementer by making a simple adder. Later in our computer we will need another adder, so we will re-use this code. Therefore, the adder code should be stored in ARM-Lab/code/0_common because it will be used in multiple stages. When used as the program counter, we will pass it a 4 because each instruction is 32-bits long (even though it is a 64-bit computer) and we want to increment to the next instruction in memory. Most machines are byte addressable, because one ASCII character (a char in c/c++) is a byte. For a machine with 32-bit instructions like we are using, that would mean that each instruction would be 4 bytes later in memory ($32/8 = 4$ bytes). Therefore,

Figure 2.1: Adder Test Simulation Results



we will be adding 4 to the program counter each time we want to increment the program counter.

An adder is very simple in Verilog. There are two inputs (the two numbers to be added) and one output (the result). All the ports are size 'WORD' because they hold 64-bit integers.

In this lab you will make your own adder module. Your adder module should be called 'adder' and should have inputs of `a_in` and `b_in`. The output should be `add_out`. HINT: this should be very easy. Verilog is a Hardware Description Language, so use Verilog to describe what you want to do. Don't make it complicated. The adder code should be stored in `ARM-Lab/code/0_common/adder.v`. You will need to create this file.

2.2 Adder Test Bench

I have provided an incomplete test bench for the adder. While I put most of the infrastructure in place, you will need to fill in the following details. I labeled these spots with comments that start with `TODO` so that you can see where you need to do your work.

1. Create an instance of the adder module
2. Fill in the details of each test case so that your simulation results and test log match the results shown in Figure 2.1 and Figure 2.2. Pay careful attention to detail to make sure your results are identical to these figures. This includes making sure that the timing matches.

2.3 Input Selection via Mux

We will also need to be able to choose between normal advancing (sequential stepping) and branching (loops, if statements, etc.). We will use a multiplexor (mux) to do this. A mux is a simple device that connects one of the inputs to the output based on how the control bit is set. If the control bit is 0 then input a is connected to the output, and if the selector is 1 then input b is connected to the output. The heading and port list is provided for you in Listing 2.1. You just need to fill in the body to make it operate like a mux.

Figure 2.2: Adder Test Log

```

***** BEGIN TEST RESULTS *****

Test Case 1
Inputs: a = 5 | b = 10
+++ Step 1: Pass: |sum| time = 7 ns | er = 15 | ar = 15 | er_bits = 64 | ar_bits = 64 +++

Test Case 2
Inputs: a = 280 | b = 1000
+++ Step 1: Pass: |sum| time = 11 ns | er = 1280 | ar = 1280 | er_bits = 64 | ar_bits = 64 +++

Test Case 3
Inputs: a = 280 | b = -1000
+++ Step 1: Pass: |sum| time = 13 ns | er = -720 | ar = -720 | er_bits = 64 | ar_bits = 64 +++

Test Case 4
Inputs: a = -280 | b = 1000
+++ Step 1: Pass: |sum| time = 18 ns | er = 720 | ar = 720 | er_bits = 64 | ar_bits = 64 +++

Test Case 5
Inputs: a = -54321 | b = 1000
+++ Step 1: Pass: |sum| time = 24 ns | er = -53321 | ar = -53321 | er_bits = 64 | ar_bits = 64 +++

Pass Count = 5
Fail Count = 0

***** END TEST RESULTS *****

```

One interesting addition in this block of code is the addition of a size parameter. Parameters allow you to pass a value into a module, making the module more flexible and reusable. In the mux code, the parameter is specified by `#(parameter SIZE=8)`. Parameters are constants, so they cannot be explicitly changed within the module. Rather, they are specified when an instance of this module is created. In the provided `mux_test.sv`, the mux module is instantiated by: `mux #(‘WORD) UUT_64`. Note that `‘WORD` is a macro (from `definitions.vh`) that defines another name for the value 64. So the test bench passes in a parameter value of 64, and the mux module uses 64 anywhere the term `SIZE` is used. If no parameter is specified when the instance is created, then the parameter will be assigned the default value. `SIZE = 8` defines the default value for the `SIZE` parameter in this module.

Note parameters are constants and cannot be changed later in the module. In this lab, we are using parameters to set the number of wires that compose the inputs and output. In our lab project, we will need some muxes to switch entire words (64 bits) like this one, but later we will also need to switch register addresses (5 bits). Rather than write two muxes, we will make one and then use the parameter to change the size when they are declared. The mux code should be stored in `ARM-Lab/code/0_common/mux.v`.

The provided testbench instantiates both a 64 bit mux and a 5-bit mux so that two different mux sizes can be tested, ensuring that the mux works now and will work when we use it later in the semester.

Listing 2.1: Verilog code to make a mux.

```
'include "definitions.vh"

module mux#(
    parameter SIZE=8)(
    input  [SIZE-1:0] a_in ,
    input  [SIZE-1:0] b_in ,
    input  control ,
    output [SIZE-1:0] mux_out
    );

    // TODO: Add body of mux here

endmodule
```

2.4 Your Assignment

You are to:

1. Create an adder module.
2. Update adder_test.sv as described above.
3. Fill in the body of the mux module in mux.v.
4. Use the provided mux_test.sv to verify that the mux works properly. Note that you cannot/should not make any changes to the test bench. The correct results are already in the test bench.
5. Produce a landscape mode PDF called Lab2_lastname.pdf that includes (in this order):
 - (a) Your name and the lab number.
 - (b) A snip of the Simulation Results for the adder_test. Make sure to show all values in decimal form and don't cut off the signal names on the left.
 - (c) The adder_test results copied and pasted from the Tcl Console. The results should show the entire log from BEGIN TEST RESULTS to END TEST RESULTS.
 - (d) A snip of the Simulation Results for the mux_test. Make sure to show all values in decimal form and don't cut off the signal names on the left.
 - (e) The mux_test results copied and pasted from the Tcl Console. The results should show the entire log from BEGIN TEST RESULTS to END TEST RESULTS.

6. Upload Lab2_lastname.pdf file to Canvas.
7. Zip up your ARM-Lab directory and submit it on Canvas as well. Please make sure that you give me the ARM-Lab directory rather than the ARM-Project directory. I do not want the project files in the ARM-Project directory. Before you submit your zip file, make sure that it contains directories like code, manual, tables, etc. I will extract your zip file and run your code against my correct testbench to verify that your code and testbench work correctly.

Lab 3

Fetch Stage

We are ready to build our fetch stage. To do this, we will make one more module, our instruction memory. Then we will make a module to assemble all of our modules together into a working fetch stage.

3.1 Instruction Memory Module

Instructions are stored in memory and are accessed by using the address where they are stored. You can think of memory like a giant hotel for our data. Each piece of data is stored in a room (memory location), which we can find by its room number (memory address). To get a piece of data stored in memory (like an instruction) we need to take its address, go to that location, and grab the data. In Verilog, a bunch of memory locations that are accessed by an address is called an array. Arrays in Verilog are declared like they are in C; the data type is specified, then the name, then the array size.

To store the instructions, we will need an array of 32-bit numbers (definitions.vh defines INSTR_LEN as 32, please use this macro), which means the data type must be `reg['INSTR_LEN-1:0]`. After the name is specified (imem in this case), we are going to use a parameter called SIZE to specify how many elements the array has: `[SIZE-1:0]`. Therefore, your array will be defined as `reg['INSTR_LEN-1:0] imem [SIZE-1:0]`.

Now we need to populate this array with instructions. Rather than populating the array element by element, we will read the instruction values in from a file called instrData.data that I have provided in the testfiles area of the ARM-Lab repository. Note that we are just initializing the array with the values from the file. The values should only be read from the file once at the beginning of the simulation. Then we will access the imem array for an instruction value. To read the file and put the contents into our imem array, we will use \$readmemb to read in hexadecimal values from instrData.data. \$readmemb could be used if we chose to format our data file in binary rather than hexadecimal. It is very important that we use a macro in definitions.vh for the name and path of the in-

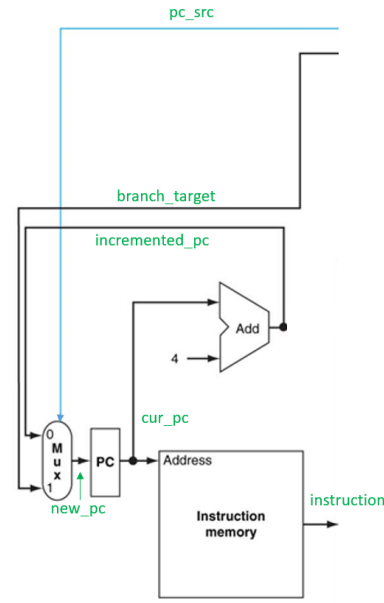
strData.data file, as this might change and it will be necessary for grading. Make sure to update the definition of IMEMFILE to point to your ARM-Lab/testfiles directory (as opposed to the potter/ARM-Lab/testfiles directory). The line of code to read the data in from the file is `$readmemh('IMEMFILE', imem);`

Once the array is populated, we can access imem and provide the instruction that corresponds to the requested address. Instructions should only be updated on the positive edge of the clock. I have provided a complete testbench for this module. I have also provided a partially complete file called instr_mem.v in code/1_fetch directory. Complete the module code to make a working instruction memory module. Test it against instr_mem.test.sv and verify that the module works as expected. The test will compare your instruction output to the instruction values that I provided in instrData.data.

3.2 Fetch Stage

Now we need to connect our modules together to make a fetch stage. The components of our instruction fetch (sometimes called ifetch or just fetch) stage are shown in Figure 3.1.

Figure 3.1: Instruction Fetch Stage.



Any wire (or reg) that comes into or goes out of the figure are input or output ports of the iFetch module. In Figure 3.1, the blue wire is a control signal and comes ultimately from the control unit, which you will build in the decode stage. Wires (or regs) that are completely contained in the figure are local to the iFetch module and are thus defined internally in the module (not an input or output). The one exception to this is the current program counter (cur_pc). While there is no reason (at this point) that it must be output from the iFetch module, you must still make it an output so that it shows up on your simulation results, helping you to keep track of the program counter for the instruction that is currently executing. And it is required to verify functionality with the testbench.

While the input and output signals are easily identified by the diagram, you must also determine the size of each signal and whether it is a wire or reg. When you look at the figure I cut from a figure in the book, note that I labeled every wire on the diagram in green. For the sake of consistency

and debugging, it is required that you use these names.

IMPORTANT NOTE: Throughout your entire project, your signal names should follow the convention of the Freescale Semiconductor Verilog guide, which states that signal names should be all lower case, with words separated by an underscore.

Once you have figured out all your connecting signals (wires and regs), you should identify the components you are going to use. We have already created the modules, so now we just need to tell Verilog to instantiate them in the iFetch module and connect them together. Create the iFetch module in a new file called iFetch.v in the code/1_fetch directory.

I have provided a complete testbench called iFetch_test.sv. Use this testbench to determine if your fetch stage is working properly. As we progress through this lab project, you will learn how critical timing is. Please look at the cur_pc value and the instruction value and verify that the instruction that was fetched is the correct instruction, according to instrData.data and the current program counter. Note that no instruction should be fetched in the first 5ns, as this is a half clock cycle and does not have a rising edge.

I have included a file called delay.v in code/0_common. It includes a module that inputs a clock signal and outputs a clock signal that is delayed by some number of ns. This will be useful when resolving timing issues. Throughout this entire lab project, the delay module should only be instantiated in the top-level module (testbench). If you need the output of the delay function in a lower-level module, add a port to the lower-level module and pass it through that port.

3.3 Your Assignment

You are to:

1. Complete the instruction memory module.
2. Verify the module by running against the provided testbench.
3. Create the iFetch stage.
4. Verify the module by running against the provided testbench.
5. Produce a landscape mode PDF called Lab3_lastname.pdf that includes (in this order):
 - (a) Your name and the lab number.
 - (b) A snip of the Simulation Results for the instr_mem_test. Please show all values in decimal except for the instructions. Please show instructions in hex.
 - (c) The instr_mem_test results copied and pasted from the Tcl Console. The results should show the entire log from BEGIN TEST RESULTS to END TEST RESULTS.

- (d) A snip of the Simulation Results for the iFetch_test. Please show all values in decimal except for the instructions. Please show instructions in hex.
 - (e) The iFetch_test results copied and pasted from the Tcl Console. The results should show the entire log from BEGIN TEST RESULTS to END TEST RESULTS.
6. Upload Lab3_lastname.pdf file to Canvas.
 7. Zip up your ARM-Lab directory and submit it on Canvas as well. Please make sure that you give me the ARM-Lab directory rather than the ARM-Project directory. I do not want the project files in the ARM-Project directory. Before you submit your zip file, make sure that it contains directories like code, manual, tables, etc. I will extract your zip file and run your code against my correct testbench to verify that your code and testbench work correctly.

Lab 4

Beginning to Decode

4.1 Instruction Decode

The next stage in the datapath is the iDecode stage. The iDecode stage evaluates the binary instructions (an output of the iFetch stage) and determines what needs to be done. There are many aspects to the iDecode stage, and some get fairly complex. But today we will begin the process of decoding an instruction by decomposing the instructions into the key parts of R-Type and D-Type instructions:

1. opcode
2. address (used only in D-Type instructions)
3. rm_num (used only in R-Type instructions)
4. rn_num
5. rd_num (though the book uses Rt for D-type instructions, we will use Rd for the last operand of D-type instructions)

To do this, you will create a new module called `instr_parse` in the `2_decode` directory. This module will simply read an input and assign appropriate output values. These outputs should be assigned using continuous assignments. The input is a 32-bit instruction. Outputs are listed for you above. Although R-type and D-type instructions have different operands, you can treat them the same for now. For instance, you can still assign an Address field on an R-type instruction, and you can still assign an Rm field on a D-type instruction. When we create the Control Module in a future lab, the control signals will drive what fields of the instruction are used and what fields are ignored. Notice how, because of the commonality of instruction format, Opcode, Rn, and Rd are all universal across these instruction types. Please remember to use the style specified in the previous lab, where all items are lower case with underscores separating them. For instance, for Rd, you should use the signal name `rd_num`. Appending num

Figure 4.1: Instruction Parse Test Output

```

***** BEGIN TEST RESULTS *****

Test Case 1
LDUR X9, [X22, #240]
+++ Step 1: Pass: |rn_num| time = 2 ns | er = 22 | ar = 22 | er_bits = 5 | ar_bits = 5 +++
+++ Step 2: Pass: |rd_num| time = 2 ns | er = 9 | ar = 9 | er_bits = 5 | ar_bits = 5 +++
+++ Step 3: Pass: |address| time = 2 ns | er = 240 | ar = 240 | er_bits = 9 | ar_bits = 9 +++
+++ Step 4: Pass: |opcode| time = 2 ns | er = 7c2 | ar = 7c2 | er_bits = 11 | ar_bits = 11 +++

Test Case 2
ADD X10, X21, X9
+++ Step 1: Pass: |rm_num| time = 12 ns | er = 9 | ar = 9 | er_bits = 5 | ar_bits = 5 +++
+++ Step 2: Pass: |rn_num| time = 12 ns | er = 21 | ar = 21 | er_bits = 5 | ar_bits = 5 +++
+++ Step 3: Pass: |rd_num| time = 12 ns | er = 10 | ar = 10 | er_bits = 5 | ar_bits = 5 +++
+++ Step 4: Pass: |opcode| time = 12 ns | er = 458 | ar = 458 | er_bits = 11 | ar_bits = 11 +++

Test Case 3
STUR X10, [X23, #64]
+++ Step 1: Pass: |rn_num| time = 22 ns | er = 23 | ar = 23 | er_bits = 5 | ar_bits = 5 +++
+++ Step 2: Pass: |rd_num| time = 22 ns | er = 10 | ar = 10 | er_bits = 5 | ar_bits = 5 +++
+++ Step 3: Pass: |address| time = 22 ns | er = 64 | ar = 64 | er_bits = 9 | ar_bits = 9 +++
+++ Step 4: Pass: |opcode| time = 22 ns | er = 7c0 | ar = 7c0 | er_bits = 11 | ar_bits = 11 +++

Pass Count = 12
Fail Count = 0

***** END TEST RESULTS *****

```

on the end of the name indicates that this is the register number, not the value from the register.

To test this module, you will need to finish `instr_parse_test.sv`. I have provided some starter code for the testbench as well as detailed comments that describe what you need to do. The testbench will feed the module with instructions. The instructions are specified in the testbench and are very similar (yet slightly different) to the instructions that were encoded in the lecture on Machine Code. You will need to update the testbench by:

1. Creating signals to be used in the testbench
2. Setting the instruction to the correct value per the instruction listed in the comments
3. Add code to verify that each output of your `instr_parse` module is correct

The final output of your testbench should match the output shown in Figure 4.1.

4.2 Register File

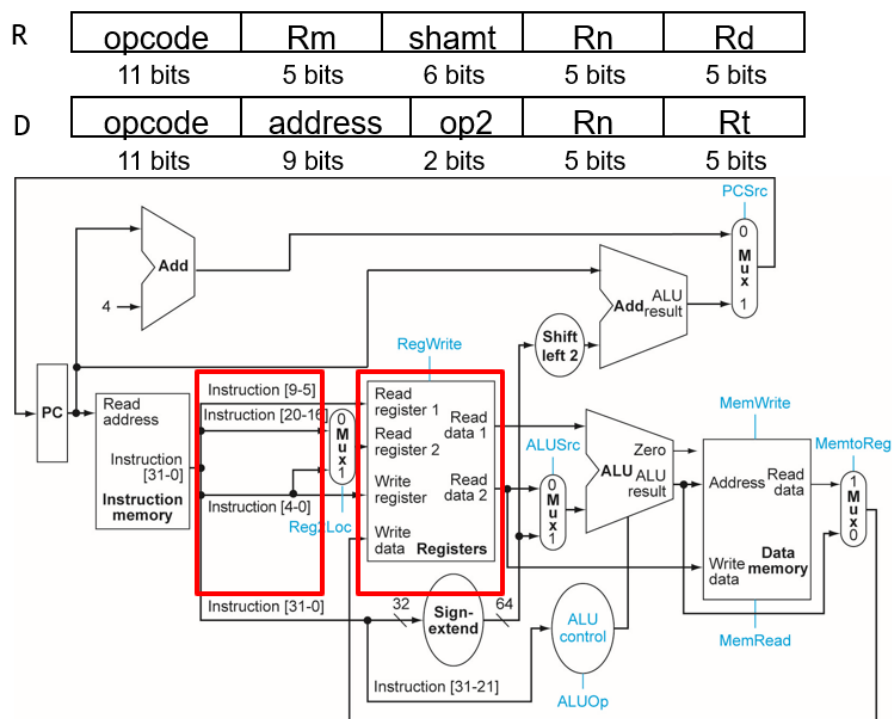
Next, we will create the register file. The register file is a piece of memory in the processor that holds the 32 register values that are used by most instructions

(X0-X31). You will create a new module called `regfile` (in a new file, `regfile.v` in `2.decode`). The `regfile` module should retrieve data from the registers on the rising edge of `read_clk` as well as write to the registers on the rising edge of `write_clk` when the `regWrite` flag is set. Two different clocks are used here because the `regfile` will be read at a different time than it is written to. The `regfile` should use a verilog `reg` array, similar to the array used in instruction memory. Since we don't currently have the ability to do loads and stores (since we don't have data memory yet), the values for the registers should be stored in a datafile, `regData.data` and copied into the array during the initial block, just like we did with the `instr_mem.v` file. `regData.data` is provided for you. The `regfile` module will have a lot of similarities to the `instr_mem` module, so I recommend reusing concepts and code from the `instr_mem` module.

Inputs to the module should include a signal called `read_clk` and a signal called `write_clk` as well as all inputs shown on the Register file in Figure 4.2. Don't forget `reg_write`. This is a control signal that determines whether data should be written to the register. Some instruction write to registers, others do not. The outputs should be the outputs of the Register file in Figure 4.2. Use names such as `read_register1`, `read_data2`, etc.

I have provided the majority of the testbench for this module, `regfile_test.sv`. It provides input values and verifies that the outputs match expected behavior. Notice how the testbench utilizes the `delay` module to create different clocks for `read_clk` and `write_clk`. The testbench is designed to test a variety of scenarios and to verify the timing aspects of this module. Your only job on the testbench is to fill in the expected result (er) values in the testbench. They are currently populated with X.

Figure 4.2: Instruction Parse and Regfile Diagrams



4.3 Your Assignment

You are to:

1. Create an `instr_parse` module as described above.
2. Update `instr_parse_test` and verify the functionality of the `instr_parse` module. For this testbench, please use hex for instructions and opcodes and unsigned decimal for all other signals.
3. Create a `regfile` module.
4. Update the `regfile_test` module as described above and verify the functionality of the `regfile` module. For this testbench, please use signed decimal for all signals.
5. Produce a landscape mode PDF called `Lab4_lastname.pdf` that includes (in this order):
 - (a) Your name and the lab number.
 - (b) A snip of the Simulation Results for the `instr_parse_test`.
 - (c) The `instr_parse_test` results copied and pasted from the Tcl Console. The results should show the entire log from BEGIN TEST RESULTS to END TEST RESULTS.
 - (d) A snip of the Simulation Results for the `regfile_test`.
 - (e) The `regfile_test` results copied and pasted from the Tcl Console. The results should show the entire log from BEGIN TEST RESULTS to END TEST RESULTS.
6. Upload `Lab4_lastname.pdf` file to Canvas.
7. Zip up your ARM-Lab directory and submit it on Canvas as well. Please make sure that you give me the ARM-Lab directory rather than the ARM-Project directory. I do not want the project files in the ARM-Project directory. Before you submit your zip file, make sure that it contains directories like code, manual, tables, etc. I will extract your zip file and run your code against my correct testbench to verify that your code and testbench work correctly.

Lab 5

Control Unit and Sign Extender

5.1 Control Unit

Next, we will create the main control unit. You will create a new module called `control` (in `control.v`) in `2_decode`. The control module should use a portion of the instruction to determine the values of all control signals to be used in our processor. These signals include:

1. `reg2_loc`
2. `uncondbranch`
3. `branch`
4. `mem_read`
5. `mem_to_reg`
6. `alu_op`
7. `mem_write`
8. `alu_src`
9. `reg_write`

The supported instructions should include:

1. `ADD`
2. `SUB`
3. `AND`

4. ORR
5. LDUR
6. STUR
7. CBZ
8. B

You will need to evaluate the incoming opcode and set the value of the control lines according to Figure 5.1. Note that a value of X in a table entry indicates that it does not matter whether the value is 0 or 1. You do not want to use X for these values, as this indicates 'undefined' to Verilog. Instead, please use 0 in place of X.

The Verilog implementation between 'if/else' statements and 'case' statements differ. 'If/else' statements will create a series of nested muxes with two inputs each, whereas a case statement will produce one large mux with many inputs. To maximize speed, we should use case statements. Note that all Verilog case statements must have a default case. But one of the challenges of this lab is dealing with opcodes of different sizes. Thankfully, Verilog has 'casex', which will only evaluate binary digits that are not labeled X. So for a CBZ instruction, you can fill in the last 3 digits with XXX and use 'casex'. The 'casex' syntax will also be very helpful in the Sign Extender. Also, please create macros in definitions.vh for the opcodes of each instruction and for the ALU Op values of each instruction type. The values in the opcode macros can include X's when necessary so that they work appropriately with the casex. For opcode macros, use the instruction type in all caps. For instance, your macro for a load instruction should be LDUR. For ALU Op macros, use the following macro names:

1. ALUOp_RTYPE
2. ALUOp_DTYPE
3. ALUOp_B
4. ALUOp_CBZ

The default case should set all control signals to 0 to "turn off" the datapath in case of an invalid opcode. The ALU Op should be R Type.

Note that you should not use a clk in this module.

5.2 Control Unit Test

The Control Unit is crucial to operation of your datapath, so it needs to be tested thoroughly. Every supported instruction (listed above) should be tested with the Control Unit Test. To facilitate testing, I have provided a testbench that includes all opcodes and an invalid opcode. You just need to fill in the 'er' values.

Figure 5.1: Control Value Table

Instruction	Reg2Loc	ALUSrc	MemtoReg	RegWrite	MemRead	MemWrite	Branch	ALUOp1	ALUOp0
R-format	0	0	0	1	0	0	0	1	0
LDUR	X	1	1	1	1	0	0	0	0
STUR	1	1	X	0	0	1	0	0	0
CBZ	1	0	X	0	0	0	1	0	1

5.3 Sign Extender

The final major component of the Decode stage is the Sign Extender. The Sign Extender should use information in the instruction to create a 64-bit output value to use as an address or branch offset. You need to extract the constant value from the instruction (this will be different for each instruction type), place that extracted value in the least significant bits of the `sign_extended_output`, then sign extend that value. For positive values, the sign extender should fill extended bits with 0s. For negative values, the sign extender should fill the extended bits with 1s. The sign extender should support extending address values from the following instructions:

1. LDUR
2. STUR
3. CBZ
4. B

For R-Type instructions that do not have a value to sign extend, please use the default case and set the value to 0.

Note that you should not use a `clk` in this module.

5.4 Sign Extender Test

The Sign Extender Test is provided for you with all 'er' values filled in. However, half of the instructions are missing. The comments indicate what the instructions should be, and you must fill in the missing instructions with the correct machine code.

5.5 Your Assignment

You are to:

1. Create a control module.
2. Update control_test.sv with the correct 'er' values and verify that your control module works properly.
3. Create a Sign Extender module.
4. Update sign_extender_test.sv to fill in the missing instructions, then verify that your sign extender module works properly.
5. Produce a landscape mode PDF called Lab5_lastname.pdf that includes (in this order):
 - (a) Your name and the lab number.
 - (b) A snip of the Simulation Results for the control module. Please show opcodes in hex and everything else in binary.
 - (c) A snip of the test results from the Tcl Console for the control module. This snip should show the entire log from BEGIN TEST RESULTS to END TEST RESULTS.
 - (d) A snip of the Simulation Results for the sign_extender module. Please show instructions in hex and everything else in signed decimal.
 - (e) A snip of the test results from the Tcl Console for the sign_extender module. This snip should show the entire log from BEGIN TEST RESULTS to END TEST RESULTS.
6. Upload Lab5_lastname.pdf file to Canvas.
7. Zip up your ARM-Lab directory and submit it on Canvas as well. I will run your code against my correct testbench to verify that your code and testbench work correctly.

Lab 6

Finishing Decode

6.1 iDecode Module

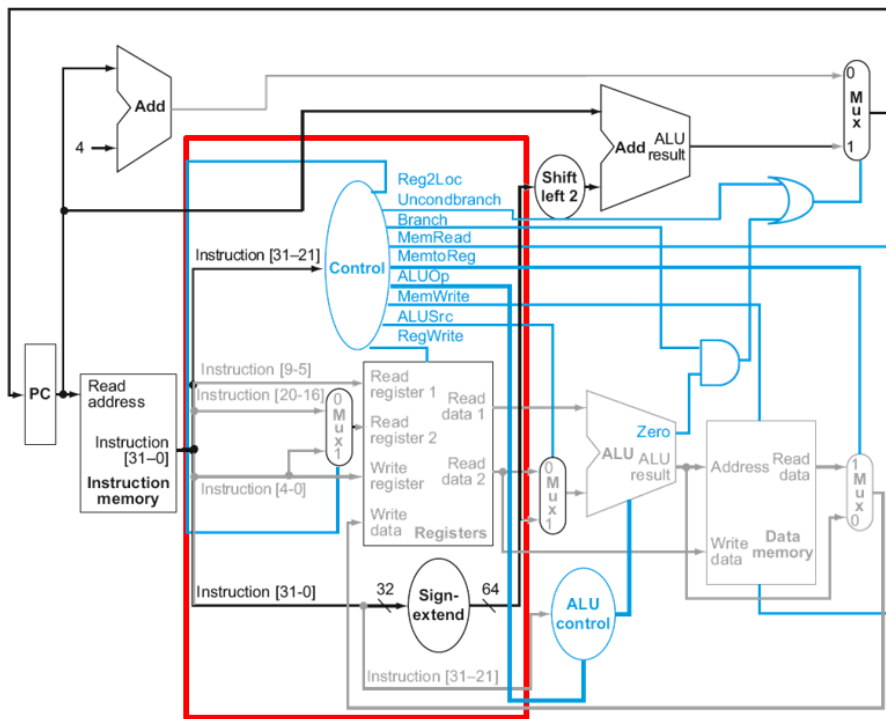
At this point, you have created all of the modules necessary to assemble the iDecode module. Now you need to create a new module called iDecode. The inputs and outputs can be determined by evaluating Figure 6.1. Any signal that crosses the boundaries of the red box is an input or output. Signals that do not cross the boundaries of the red box are signals that are internal to the iDecode module and should be declared internally in iDecode, with the exception of `reg_write` and `reg2_loc`. Please make sure to label signals consistently with lower case letters with words separated by underscores. For example, `read_data1`, `write_data`, `alu_src`.

6.2 iDecode Testbench

To verify that your iDecode module works correctly, you must first fill in the Expected Results Table that has been provided in the tables directory. All rows and columns have been defined, but you must fill the expected values into each cell. Note that, at this point, you cannot fill in the rows below `write_data`. If a particular data item is not relevant for a particular instruction (for instance, `sign_extended_output` on an R-type instruction), then put N/A in that cell. It is important that we identify which signals are applicable and which signals are not applicable. Please do not use N/A for any control signals, as we always want to make sure these are set correctly.

To ensure that we are testing each case correctly and consistently, please update `regData.data` to reflect the following values and use these values when making filling in your Expected Results Table. Note that the instructions execute, so if X9 is updated by the first instruction, then you should use the updated value of X9 in the second instruction, etc. However, we will not be branching at this time, so the instructions will run in order.

Figure 6.1: Expected Results



1. $X19 = 10$
2. $X20 = 30$
3. $X21 = 0$
4. $X22 = 16$

Now update the provided testbench for iDecode to provide the inputs for iDecode and verifying the outputs of iDecode. For the instruction input, use the instructions from your Expected Results Table. For the outputs, use your Expected Results Table to fill in the 'er' values for each instruction. For instructions that update the register file, use the test bench to provide the correct value to the write_data input, since we do not yet have an ALU to do the calculations. Also, please use your test bench to provide a value of 20 to X9 in the first command (LDUR).

In the testbench, when there is a value that is N/A in your Expected Results Table, please comment out the verify function call for that signal, since we do not have a real expected result to compare it to. Doing this, my testbench ends up having 119 test cases.

6.3 Your Assignment

You are to:

1. Integrate all individual modules into the iDecode module.
2. Update your Expected Results Table to include values for all iDecode inputs and outputs.
3. Update iDecode_test.sv with values from the Expected Results Table. Note that you do not need to add any test cases, change timing, etc.
4. Verify that your simulation results match your expected results.
5. Rather than writing a lab report, please produce a landscape mode PDF file called Lab6_lastname.pdf that includes (in this order):
 - (a) Your name and the lab number.
 - (b) A snip of your completed Expected Results Table.
 - (c) A snip of the Simulation Results for the iDecode module. Please show instructions in hex, opcodes and control signals in binary and everything else in signed decimal.
 - (d) A snip of the test results from the Tcl Console for the iDecode module. This snip should show the entire log from BEGIN TEST RESULTS to END TEST RESULTS.
6. Upload Lab6_lastname.pdf file to Canvas.

7. Zip up your ARM-Lab directory and submit it on Canvas as well. I will run your code against my correct testbench to verify that your code and testbench work correctly.

Lab 7

Integrating Fetch and Decode

This lab will integrate the iFetch module with the iDecode module to make a single testbench that will verify that they function properly together. The result will be a system that automatically fetches instructions and then decodes the fetched instruction. To do this, you will update the provided `fd_integration.sv` to include an instance of both the iFetch module and the iDecode module. The instruction output from the fetch module should be connected to the instruction input in the decode module.

You will also need to thoroughly test the outputs of each stage. Refer back to your `iFetch_test.sv` and `iDecode_test.sv` files to identify what you need to test. You will need to verify every item that you verified in each of those previous tests. We will be using the instructions and expected values from the Expected Results Table in the iDecode lab. Note that this time, rather than providing the instructions in the testbench, the iFetch module should provide the instructions. This means that you will need to update `instrData.data` with the 10 instructions from the Expected Results Table. Make sure to also verify the `cur_pc` and instruction outputs of the iFetch module. Overall, the tests should have 139 steps. 119 come from the iDecode signals and 20 come from the iFetch signals.

7.1 Your Assignment

You are to:

1. Update the provided `fd_integration.sv` by instantiating iFetch and iDecode modules. Also, update `fd_integration.sv` with values from the Expected Results Table and with expected outputs from iFetch.
2. Verify that your simulation results match your expected results.

3. Rather than writing a lab report, please produce a landscape mode PDF file called Lab7_lastname.pdf that includes (in this order):
 - (a) Your name and the lab number.
 - (b) A snip of your completed Expected Results Table.
 - (c) A snip of the Simulation Results for the fd_integration test. Please show instructions in hex, opcodes and control signals in binary and everything else in signed decimal.
 - (d) Copy and paste the entire log from BEGIN TEST RESULTS to END TEST RESULTS into your file. The results have gotten too long to use the snipping tool.
4. Upload Lab7_lastname.pdf file to Canvas.
5. Zip up your ARM-Lab directory and submit it on Canvas as well. I will run your code against my correct testbench to verify that your code and testbench work correctly.

Lab 8

ALU and ALU Control

The goal of this lab is to build the ALU and ALU Control modules. These modules can be seen in the context of the datapath by viewing the datapath diagram from a previous lab, Figure 6.1. The ALU performs the arithmetic and logical operations, and the ALU Control module determines which operation should be performed by the ALU. During this lab, we will not put these modules together. Instead, they will be tested independently of each other.

8.1 ALU

The ALU has three inputs:

1. `a_in` - the first input operand
2. `b_in` - the second input operand
3. `alu_control` - control signal used to tell the ALU what operation to perform

The ALU has two outputs:

1. `alu_result` - the result of the arithmetic/logic operation
2. `zero` - a flag indicating whether `alu_result` is zero

Figure 8.1 identifies the operation that corresponds to the `alu_ctrl` value. Note that we will not be implementing NOR, even though it is listed in the table. You should use a case statement to evaluate the `alu_control` bits to determine which ALU operation to perform. For each operation, you do not have to do anything fancy. You just need to use the math capability that verilog provides to make the calculation. To make the code readable, you must give the the ALU control bits names in the `definitions.vh` file, and you should use these in your cases. Please use the following names for your macros in `defintions.vh`:

1. `ALU_AND`

Figure 8.1: ALU Control Values

ALU control lines	Function
0000	AND
0001	OR
0010	add
0110	subtract
0111	pass input b
1100	NOR

2. ALU_ORR
3. ALU_ADD
4. ALU_SUB
5. ALU_PASS

Also, don't forget to make a default case, which is needed to actually wire this up. Please use AND for the default case.

One last thing to note is the generation of the zero flag. The zero flag is determined by simply evaluating the `alu_result`. There are several ways to handle this, but this is the easiest way to handle it. In Verilog (like C), the statement `(y == 0)` is an operation with a boolean output. You can thus say `x = (y == 0);` to assign `x` to be the boolean value that `(y == 0)` produced. The statement `x = (y == 0);` is realizable as a digital comparator with `y` and 0 as inputs and `x` as the single bit output. Please note that you will not have signals called `x` and `y`. I just used these to explain the concept.

8.2 ALU Control

The ALU Control module has two inputs:

1. `alu_op` - 2 bit signal giving incomplete information about what ALU operation should be performed
2. `opcode` - the opcode of the instruction

The ALU Control module has one output:

1. `alu_control` - control signal used to tell the ALU what operation to perform

You have already defined the `ALUOp` macros in `definitions.vh` when you created the control module. You have also already defined the ALU Control macros in `definitions.vh` in the previous section. And you have macros defined for

Figure 8.2: ALU Op Table

ALUOp		Opcode field											Operation
ALUOp1	ALUOp0	I[31]	I[30]	I[29]	I[28]	I[27]	I[26]	I[25]	I[24]	I[23]	I[22]	I[21]	
0	0	X	X	X	X	X	X	X	X	X	X	X	0010
X	1	X	X	X	X	X	X	X	X	X	X	X	0111
1	X	1	0	0	0	1	0	1	1	0	0	0	0010
1	X	1	1	0	0	1	0	1	1	0	0	0	0110
1	X	1	0	0	0	1	0	1	0	0	0	0	0000
1	X	1	0	1	0	1	0	1	0	0	0	0	0001

instruction opcodes. Please use these macros when creating the ALU Control module.

Figure 8.1 identifies the operation that corresponds to the `alu_ctrl` value. This module can be implemented multiple ways, but we want to use the most efficient way possible. The most efficient method is to use a `casex` statement to evaluate the `alu_op` signal, utilizing the information in Figure 8.2. `alu_op` will provide all information necessary for D-Type and CB-Type instructions. For R-Type instructions, you should use the bits of the opcode to set the `alu_ctrl` value. When evaluating Figure 8.2 for R-Type instructions, certain bits of the instruction correspond to `alu_ctrl` bits. The magic decoder ring is listed below. Please note that the numbering system for the instruction bits includes the entire instruction, whereas you will just have the opcode to work with. Therefore, you will need to adjust the numbering to account for your 11-bit opcode.

1. ALU Ctrl [3] = 0
2. ALU Ctrl [2] = I[30]
3. ALU Ctrl [1] = I[24]
4. ALU Ctrl [0] = I[29]

Please use AND as the default case. For the B instruction, the ALU Op should be 00 from your control module. Since 00 is also the ALU Op for LDUR and STUR (which have ALU Control value of ADD), you can just allow the B instruction to also set ALU Control to ADD.

8.3 Your Assignment

You are to:

1. Create the ALU module.
2. Verify that your ALU module works correctly by running it against the provided test bench.
3. Create an ALU Control module.

4. Verify that your ALU module works correctly by running it against the provided test bench.
5. Rather than writing a lab report, please produce a landscape mode PDF file called Lab8_lastname.pdf that includes (in this order):
 - (a) Your name and the lab number.
 - (b) A snip of the Simulation Results for the each test bench. Please show instructions in hex, opcodes and control signals in binary and everything else in signed decimal.
 - (c) For each test bench, copy and paste the entire log from BEGIN TEST RESULTS to END TEST RESULTS into your file.
6. Upload Lab8_lastname.pdf file to Canvas.
7. Zip up your ARM-Lab directory and submit it on Canvas as well. I will run your code against my correct testbench to verify that your code and testbench work correctly.

Lab 9

Execute Stage

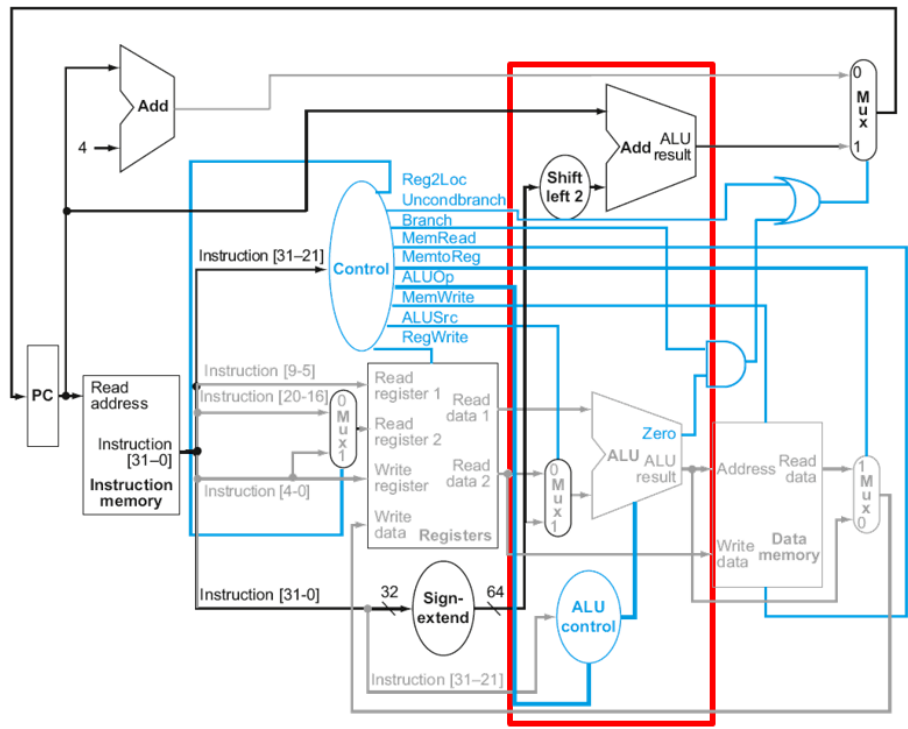
9.1 Execute

In the last lab, you created the ALU and ALU Control modules. Now we will finish the iExecute stage by creating a new module called iExecute. The iExecute module should consist of everything shown in the red box in Figure 9.1. To implement the iExecute stage, you will need to create iExecute.v and add the following components to it:

1. ALU module from the previous lab
2. ALU Control module from the previous lab
3. Mux to select the source of the second input into the ALU. You can reuse your mux that you created in the iFetch stage.
4. Shifter to left shift the sign extended branch address offset. You should NOT create a new module for this. Instead, you should add code directly in iExecute.v since the left shift can be accomplished with a single line of code.
5. Adder to add the branch address offset to the current PC. You can reuse your adder that you created in the iFetch stage. Please note the adder is near the top of the datapath diagram and it has an output named ALU result on the diagram. In fact, this is not a full ALU, but just an adder. And our output signal will be named branch_target, as this is what it really is (rather than naming it ALU_Result).

The testbench will run the 10 instructions in your Expected Results Table. Therefore, your next step is to update your Expected Results Table by updating the 3 outputs of the iExecute stage. You also need to update 1 input to the table, PC. The PC value should start at 0 and increment by 4 for each instruction. Populate the iExecute output rows with expected values for each instruction.

Figure 9.1: Execute Stage



All of the inputs that you need to determine the expected result are in the table. Make sure to use N/A if the signal is not applicable for a given instruction.

Finally, you need to complete the testbench. I have provided the bulk of the testbench. The only updates that you need to make to the testbench are:

1. I have X for all er values right now. Please update these to the correct values per your Expected Results Table.
2. I have a verify function call for all 3 outputs for every instruction. If an output is N/A for a particular instruction, remove or comment out that verify. There should be a total of 20 test cases.

9.2 Your Assignment

You are to:

1. Complete the iExecute module
2. Update your Expected Results Table with the outputs from the iExecute stage. Also add a PC value.
3. Update iExecute_test.sv
4. Verify that your simulation results match your expected results.
5. Rather than writing a lab report, please produce a landscape mode PDF file called Lab9_lastname.pdf that includes (in this order):
 - (a) Your name and the lab number.
 - (b) A snip of your completed Expected Results Table.
 - (c) A snip of the Simulation Results for the iExecute test. Please show instructions in hex, opcodes and control signals in binary and everything else in signed decimal.
 - (d) Copy and paste the entire log from BEGIN TEST RESULTS to END TEST RESULTS into your file. The results have gotten too long to use the snipping tool.
6. Upload Lab9_lastname.pdf file to Canvas.
7. Zip up your ARM-Lab directory and submit it on Canvas as well. I will run your code against my correct testbench to verify that your code and testbench work correctly.

Lab 10

Integrating Fetch and Decode

10.1 Integration

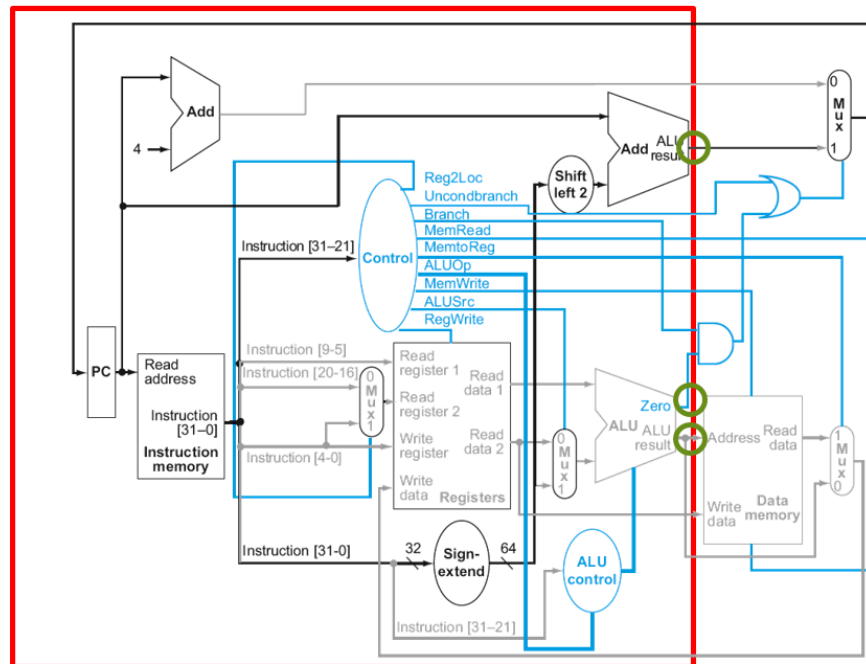
We now have working Fetch, Decode, and Execute modules. Now it is time to put them together to produce a system that can:

1. Update the program counter
2. Read the appropriate instruction from the instruction datafile
3. Read the correct registers
4. Update all control lines
5. Sign extend address data
6. Calculate Branch Target Addresses
7. Provide a zero bit for conditional branch instructions
8. Produce ALU results for R-Type and D-Type instructions

Once we can do all of this, we will be ready for the iMemory stage. We currently have the fetch and decode module integrated into `fd_integration.sv`. We also have a working execute module. Today we need to integrate the execute module with the fetch and decode stages. The new testbench should be called `fde_integration.sv`. Please reuse the instructions from your Expected Results Table that you used when you integrated fetch and decode. Once integrated, you should be able to produce a simulation that includes the outputs of fetch and decode as well as 3 new outputs from execute:

1. Branch Target

Figure 10.1: Execute Stage



2. ALU Result
3. Zero

These three new outputs are marked on Figure Figure 10.1. To verify these outputs, you should have already updated the Expected Results Table to include these three outputs. To test this integration, I recommend copying the contents of `fd.integration.sv` into `fde.integration.sv` and then modifying from there, adding your `execute` module, `cr` values, `verify` statements, etc that are necessary to verify the integration of the `execute` module. `fd.integration` had 139 test cases, so `fde.integration` should have 159 test cases.

10.2 Your Assignment

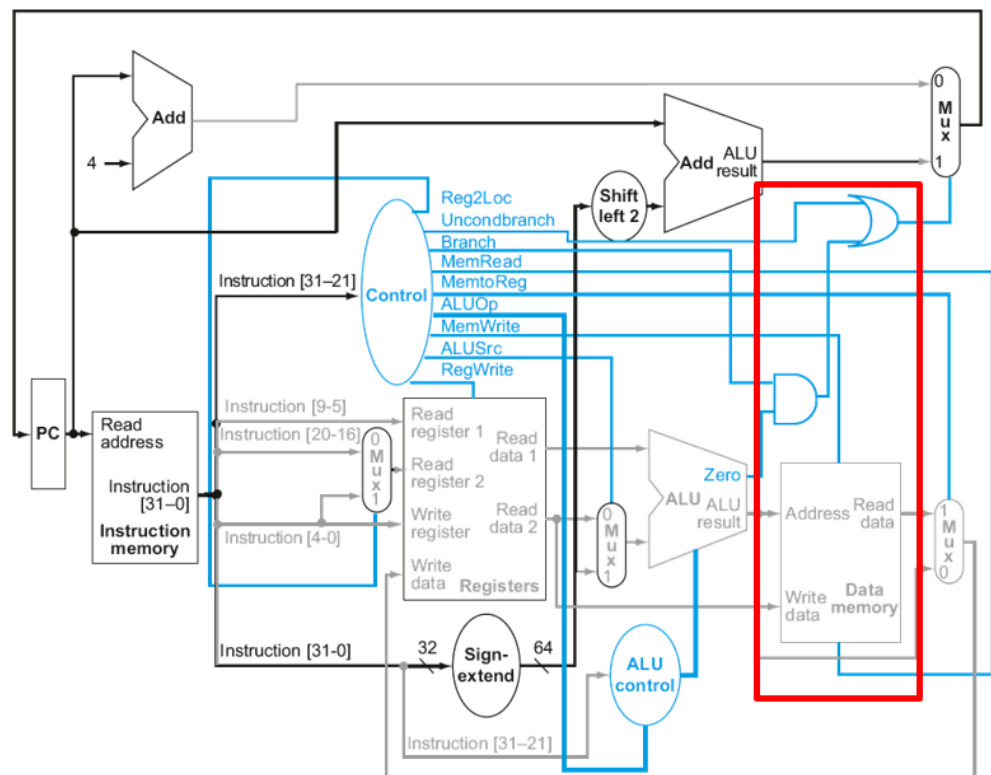
You are to:

1. Create `fde_integration.sv`
2. Verify that your simulation results match your expected results.

3. Rather than writing a lab report, please produce a landscape mode PDF file called Lab10_lastname.pdf that includes (in this order):
 - (a) Your name and the lab number.
 - (b) A snip of your completed Expected Results Table.
 - (c) A snip of the Simulation Results for the iExecute test. Please show instructions in hex, opcodes and control signals in binary and everything else in signed decimal.
 - (d) Copy and paste the entire log from BEGIN TEST RESULTS to END TEST RESULTS into your file. The results have gotten too long to use the snipping tool.
4. Upload Lab10_lastname.pdf file to Canvas.
5. Zip up your ARM-Lab directory and submit it on Canvas as well. I will run your code against my correct testbench to verify that your code and testbench work correctly.

Lab 11

Memory



11.1 Memory Stage

Today we will create the iMemory stage of our processor. This stage contains the data memory for the system that we use for load and store commands. It also contains the logic gates used to produce the pc_src signal that is used in the iFetch stage. Note that although the diagram shows the pc_src mux on the right side of the diagram, the mux is actually already implemented in the iFetch stage and belongs in the iFetch stage.

Because this module is so simple, we will put all of the code for the module in a new file called iMemory.v. This file will read and write data memory as well as determine the branch result.

11.2 Data Memory

The data memory will be similar to the register file memory, with two primary changes:

1. reading is now conditional on the MemRead control wire being high. If the MemRead flag is not high, then the read_data output should be set to Z (high impedance).
2. writing is now permissible if the MemWrite control wire is high.

In iMemory.v, copy the contents of register_file.v and modify it to meet the needs of the data memory module. I have provided a new data file, ramData.data that contains the initial values to be read into memory.

11.3 Branch Resolution

We now have all the information necessary to decide if the computer should branch or not. We have the signal 'branch' to tell us if it is a conditional branch command, and we have 'zero' to tell us if the condition was met. Both branch and zero must be true so we will combine them with an 'and' gate.

We also need an 'or' gate to 'or' together the output of the 'and' gate (above) and the uncondBranch control signal. These gates can be included directly in your iMemory module. While these are two separate gates on the Datapath Diagram, the code can be written as a single line of code (or you can write it as two lines, either way is fine).

11.4 Test Bench

For this lab, we will just have a single test bench that will test the entire iMemory stage. The provided test bench is complete except that the cr values are not filled in. Please fill in these values with the correct results that you expect to get.

11.5 Your Assignment

You are to:

1. Create a new module called iMemory as described above.
2. Update the testbench cr values to verify that the iMemory stage works properly.
3. Rather than writing a lab report, please produce a landscape mode PDF file called Lab11_lastname.pdf that includes (in this order):
 - (a) Your name and the lab number.
 - (b) A snip of the Simulation Results for the iMemory test. Please show instructions in hex, opcodes and control signals in binary and everything else in signed decimal.
 - (c) Copy and paste the entire log from BEGIN TEST RESULTS to END TEST RESULTS into your file. The results have gotten too long to use the snipping tool.
4. Upload Lab11_lastname.pdf file to Canvas.
5. Zip up your ARM-Lab directory and submit it on Canvas as well. I will run your code against my correct testbench to verify that your code and testbench work correctly.

Lab 12

Write Back

12.1 Mux

This stage consists on only one item, a mux to select between the output of memory and the output of the ALU. The control is the MemtoReg control line, see Fig 12.1. Since the mux has already been tested it does not need a testbench. The stage thus has only 3 inputs (2 data and 1 control) and one output, the result.

12.2 Datapath

You are ready to assemble the full non-pipelined datapath shown in Fig 12.2. To do this, you will need to combine all 5 stages into datapath.sv. datapath.sv is a new top-level file, so it is the testbench. Verify operation of your datapath by running your set of instructions in instrData.data and testing the output. These instructions are the 10 instructions from the Expected Results Table.

Now that we have a writeback stage, we do not need to set write.data in the initial section of datapath.sv. Rather, you should connect write.data from the WriteBack stage to the Decode stage. Because we now have a memory stage, we should no longer need to set pc.src in the initial section of datapath.sv. However, our test instructions are not meant to run like a program and would yield strange results. So for right now, we want to keep pc.src hard-coded to 0 in datapath.sv. To do this, we will use a new reg called pc.src.tmp and set it to 0. Then we will use pc.src.tmp as the input to the iFetch module. pc.src will be a wire that is an output of the iMemory module. In part 2 of this lab, we will get rid of pc.src.tmp and connect pc.src from iMemory to iFetch.

12.3 Your Assignment Part 1

You are to:

Figure 12.1: Write Back





Figure 12.2: Full Non-Pipelined Datapath



1. Update your Expected Results Table to include the iWriteBack stage.
2. Create the iWriteback stage consisting of one Mux.
3. Integrate all five stages into the file datapath.sv.
4. Run simulations to verify that your results match your Expected Results table.
5. There will be one final test to add before we submit it. However, please save copies of the following:
 - (a) A snip of the Simulation Results for datapath.sv. Please show instructions in hex, opcodes and control signals in binary and everything else in signed decimal.
 - (b) Copy and paste the entire log from BEGIN TEST RESULTS to END TEST RESULTS into your file. The results have gotten too long to use the snipping tool.

12.4 Division

To fully verify your LEGv8 processor, we are going to write a program that can divide two numbers, run that program and verify that we get the correct result. The program that we will run is shown in the file division.c, shown below. At the top of that file, I've listed some rules about the contents of the regData file and the ramData file. In previous semesters, I have made students translate this C code into assembly, then translate that assembly into machine code, and put that machine code in instrData. This semester, to lighten your load, I am providing 2 new files:

1. division_assembly.txt
2. instrData_division.data

While I am providing these files, please open them up and analyze the contents. To utilize these files, you need to copy them into your testfiles directory and update definitions.vh to point to instrData_division.data. instrData_division contains the machine code for the C code listed below. The assembly translation is also shown below. Also, you should create a new file called regData_division.data. You must set all register values in regData_division.data to 0 except for A22, which should be set to 24. This value is the base address of the array A. This is the only non-zero value allowed in regData. You should also create ramData_division. It should have all values set to 0 except for the values of A[0], A[1], A[2], and A[3]. Please set the dividend to 56 and the divisor to 8. The quotient is initially set to 0. A[3] is set to 1. This is necessary because we do not have immediate instructions in our processor.

To verify operation, you should create a final test bench (division.sv) that is the same as datapath.sv except that it allows pc_src to be set by the iMemory

module rather than keeping it hard-coded to 0. These instructions are meant to run as a program and correct branch decisions are vital to its operation. The test bench verification is very simple. It has one test case, which verifies the value that is being stored by the final instruction. Given that we are calculating $56/8$, the result should be 7. Rather than verifying hundreds of steps along the way (which we've already done), we will very simply check to see if the division works by checking the value of `read_data2` in the final instruction (STUR). Feel free to modify `ramData_division.data` to use a different dividend and divisor...but make sure that they are divisible. $57/8$ (or anything like that) will not work properly, as it will never break out of the loop.

Listing 12.1: C code for doing simple division.

```
// Rules:
// The register file must be set to all 0s except that you
// can store the base address of A in X22. Base address of A
// should be non-zero. Everything else must be loaded from memory.
// This is a simple division problem where having a remainder will
// break it, so choose an x that is divisible by y.

int x = A[0]; // dividend
int y = A[1]; // divisor
int z = A[2]; // quotient, should initially be set to 0
int one = A[3];

while(x != 0)
{
    x = x - y;
    z = z + one;
}

A[2] = z;
```

Listing 12.2: Assembly code for doing simple division.

```
LDUR X9, [X22, #0]
LDUR X10, [X22, #8]
LDUR X11, [X22, #16]
LDUR X11, [X22, #24]
CBZ X9, 4
SUB X9, X9, X10
ADD X11, X11, X12
B -3
STUR X11, [X22, #16]
```

12.5 Your Assignment Part 2

You are to:

1. Create division.sv.
2. Run the simulation, analyze the results and verify that the division works correctly.
3. Rather than writing a lab report, please produce a landscape mode PDF file called Lab12_lastname.pdf that includes (in this order):
 - (a) Your name and the lab number.
 - (b) A snip of the Simulation Results for both the datapath.sv test and division.sv. Please show instructions in hex, opcodes and control signals in binary and everything else in signed decimal.
 - (c) Copy and paste the entire log for both datapath.sv and division.sv from BEGIN TEST RESULTS to END TEST RESULTS into your file. The results have gotten too long to use the snipping tool.
4. Upload Lab12_lastname.pdf file to Canvas.
5. Zip up your ARM-Lab directory and submit it on Canvas as well. I will run your code against my correct testbench to verify that your code and testbench work correctly.

CONGRATULATIONS! YOU'VE JUST BUILT A SIMPLE ARMv8 PROCESSOR!!!!

Lab 13

Pipeline Fetch and Decode

13.1 Pipeline Fetch and Decode

Now that we have a working single-cycle datapath, we are going to break it apart and start pipelining it. Make sure to keep a copy of your working single cycle datapath for reference. To begin the pipelining process, we will first pipeline the iFetch and iDecode stages.

The advantage of pipelining includes the ability to reduce the clock cycle time by doing only one stage at a time. We have been using a 10ns clock cycle. Now that we are pipelining, let's start conservative and go down to a 4ns clock cycle. Please keep it at 4ns for now to maintain consistency across the class.

This will affect your timing and will require you to update the clocks in your datapath. The basic clk for each stage should be clk (no delays). You will still need delayed clocks for items that are delayed within a stage such as read_clk and write_clk. The first goal for your pipeline is to get a series of instructions to execute in pipeline form, where a new instruction is fetched while the previous instruction is being decoded. You should verify this by examining the Read_Register values, Read_Data values, control signal values, etc. To test this, use your set of instructions that corresponds to your Expected Results table. You should not have to make a lot of updates, as the instr_parse module was made partially for the purpose of buffering data, making pipelining easier. For right now, just focus on the Fetch and Decode stages.

13.2 Pipeline Buffering and Forwarding

The next step in developing your pipeline is to add the stage buffers that we have seen in the textbook. The idea here is that every input to a stage must be buffered. The code for this is not complex, but the planning and thought process can be complex because some signals need to be passed through to future stages. Remember, you should not have signals that go from Decode to Memory, for example. If you do, the data will be incorrect due to timing issues. To aid your

efforts, I require that you complete the PipelineAnalysisTemplate.xlsx. This allows you to figure out what signals need to be passed between stages and develops a naming convention that will help you when you start implementing.

In the spreadsheet, you should start with my template, which has many (not all) of the signals that are in your datapath. Look at your current connections in datapath.v. Wherever you see an output from a module, put that in the spreadsheet as a source. Wherever you see an input to a module, put that in the spreadsheet as a destination. In many cases, you will see that there is space in the spreadsheet between the source and destination. This means that the signal must be buffered and forwarded from one stage to the next. Fill in the space between the source and destination with signals to be forwarded. Please use the naming convention that I show on the spreadsheet....this might be the most crucial part. The signal names should match your current signal names, but they should have a suffix added to them that indicates which stage outputted the data. For instance, when mem_to_reg was buffered in iExecute and forwarded to iMemory, the signal should be name mem_to_reg_ie because iExecute outputs the data.

13.3 Your Assignment

You are to:

1. Save a copy of your non-pipelined datapath.v.
2. Comment out Execute, Memory, and WriteBack stages from datapath.v.
3. Pipeline iFetch and iDecode stages as described above.
4. Test with your set of instructions from your Expected Results table.
5. Verify that stages are processing the correct instructions at the correct times.
6. Complete PipelineAnalysisTemplate.xlsx
7. There is no lab report today.

Lab 14

Pipelining without Branching or Forwarding

14.1 Overview

Now that we have pipelined the fetch and decode stages, we can add the register buffers between each stage and get a simple pipeline working. This pipeline will not include any data forwarding or branch prediction. We will handle data hazards by using assembly code with appropriately placed nop commands. For now, we will avoid control hazards by keeping pc.src set to 0, which will keep the system from branching (even when you run a branch instruction).

Only after you have completed the spreadsheet, then you can move on to writing Verilog code. You will want to update datapath.v to use the names from the spreadsheet. This will include adding a lot of new signals. Then you will want to update your modules to account for these new signals. Use your spreadsheet to make the following updates:

1. Add input ports to stage modules as needed. This will be particularly relevant for signals that are passed through. You should not use the `_id`, `_ie`, etc naming convention within the module. You should either use no suffix or the suffix `_in` or `_out`.
2. Update all modules to buffer all inputs into the module. This will keep inputs from changing while they are being used. On the positive edge of the clock, you should use procedural assignments to copy all inputs into another register (reg). This register can either be an output reg (when appropriate) or a local reg (if the signal does not need to be output). Then do all of the module's processing on the buffered regs. Try to change as few names as possible within the modules. Use names that keep you from having to update your signal names to the lower level modules. In modules that do not currently have a clock, you will need to add the clock.

3. Use the information in the spreadsheet to add output ports to stage modules as needed.

Again, try to change as few names as possible within the modules. You want to use the stage-specific names (ie, etc) in datapath.v rather than in the modules, to the extent possible.

Use the instructions from your Expected Results Table to test your pipeline. Insert NOP instructions where appropriate in your instrData.data. To insert a NOP instruction, add a 32-bit line with all zeros.

14.2 Your Assignment

You are to:

1. Implement the spreadsheet in your datapath.v
2. Update your modules to buffer the appropriate values into registers
3. Use your expected results table instructions to test the pipeline and correct any issues
4. Once your pipeline is working, go back and try to reduce the cycle time as much as possible (without having to re-architect your solution).
5. Submit a lab report using the PipelineLabWriteup format. This is a new format that has been uploaded to my Git repository.